
Linux Sh Documentation

The kernel development community

Jun 10, 2024

CONTENTS

1	DeviceTree Booting	3
2	Adding a new board to LinuxSH	5
3	Notes on register bank usage in the kernel	11
4	Memory Management	13
5	Machine Specific Interfaces	15
6	Busses	17

Author

Paul Mundt

DEVICETREE BOOTING

Device-tree compatible SH bootloaders are expected to provide the physical address of the device tree blob in r4. Since legacy bootloaders did not guarantee any particular initial register state, kernels built to interoperate with old bootloaders must either use a builtin DTB or select a legacy board option (something other than `CONFIG_SH_DEVICE_TREE`) that does not use device tree. Support for the latter is being phased out in favor of device tree.

ADDING A NEW BOARD TO LINUXSH

Paul Mundt <lethal@linux-sh.org>

This document attempts to outline what steps are necessary to add support for new boards to the LinuxSH port under the new 2.5 and 2.6 kernels. This also attempts to outline some of the noticeable changes between the 2.4 and the 2.5/2.6 SH backend.

2.1 1. New Directory Structure

The first thing to note is the new directory structure. Under 2.4, most of the board-specific code (with the exception of stboards) ended up in arch/sh/kernel/ directly, with board-specific headers ending up in include/asm-sh/. For the new kernel, things are broken out by board type, companion chip type, and CPU type. Looking at a tree view of this directory hierarchy looks like the following:

Board-specific code:

```
.
|-- arch
|   |-- sh
|   |   |-- boards
|   |   |   |-- adx
|   |   |   |   |-- board-specific files
|   |   |   |-- bigsur
|   |   |   |   |-- board-specific files
|   |   |   ... more boards here ...
|-- include
|   |-- asm-sh
|   |   |-- adx
|   |   |   |-- board-specific headers
|   |   |-- bigsur
|   |   |   |-- board-specific headers
|   |   .. more boards here ...
```

Next, for companion chips:

```
.
|-- arch
|   |-- sh
|   |   |-- cchips
|   |   |   |-- hd6446x
|   |   |   |   |-- hd64461
|   |   |   |       |-- cchip-specific files
```

...and so on. Headers for the companion chips are treated the same way as board-specific headers. Thus, `include/asm-sh/hd64461` is home to all of the `hd64461`-specific headers.

Finally, CPU family support is also abstracted:

```
.
|-- arch
|   |-- sh
|   |   |-- kernel
|   |   |   |-- cpu
|   |   |   |   |-- sh2
|   |   |   |   |   |-- SH-2 generic files
|   |   |   |   |-- sh3
|   |   |   |   |   |-- SH-3 generic files
|   |   |   |   |-- sh4
|   |   |   |       |-- SH-4 generic files
|   |   |-- mm
|   |       |-- This is also broken out per CPU family, so each
|   |       |   ↳ family can have their own set of cache/tlb functions.
|-- include
|   |-- asm-sh
|   |   |-- cpu-sh2
|   |   |   |-- SH-2 specific headers
|   |   |-- cpu-sh3
|   |   |   |-- SH-3 specific headers
|   |   |-- cpu-sh4
|   |       |-- SH-4 specific headers
```

It should be noted that CPU subtypes are not abstracted. Thus, these still need to be dealt with by the CPU family specific code.

2.2 2. Adding a New Board

The first thing to determine is whether the board you are adding will be isolated, or whether it will be part of a family of boards that can mostly share the same board-specific code with minor differences.

In the first case, this is just a matter of making a directory for your board in `arch/sh/boards/` and adding rules to hook your board in with the build system (more on this in the next section). However, for board families it makes more sense to have a common top-level `arch/sh/boards/` directory and then populate that with sub-directories for each member of the family. Both the Solution Engine and the `hp6xx` boards are an example of this.

After you have setup your new `arch/sh/boards/` directory, remember that you should also add a directory in `include/asm-sh/` for headers localized to this board (if there are going to be more than one). In order to interoperate seamlessly with the build system, it's best to have this directory the same as the `arch/sh/boards/` directory name, though if your board is again part of a family, the build system has ways of dealing with this (via `incdir-y` overloading), and you can feel free to name the directory after the family member itself.

There are a few things that each board is required to have, both in the `arch/sh/boards` and the `include/asm-sh/` hierarchy. In order to better explain this, we use some examples for adding an imaginary board. For setup code, we're required at the very least to provide definitions for `get_system_type()` and `platform_setup()`. For our imaginary board, this might look something like:

```
/*
 * arch/sh/boards/vapor/setup.c - Setup code for imaginary board
 */
#include <linux/init.h>

const char *get_system_type(void)
{
    return "FooTech Vaporboard";
}

int __init platform_setup(void)
{
    /*
     * If our hardware actually existed, we would do real
     * setup here. Though it's also sane to leave this empty
     * if there's no real init work that has to be done for
     * this board.
     */

    /* Start-up imaginary PCI ... */

    /* And whatever else ... */

    return 0;
}
```

Our new imaginary board will also have to tie into the machvec in order for it to be of any use.

machvec functions fall into a number of categories:

- I/O functions to IO memory (inb etc) and PCI/main memory (readb etc).
- I/O mapping functions (ioport_map, ioport_unmap, etc).
- a ‘heartbeat’ function.
- PCI and IRQ initialization routines.
- Consistent allocators (for boards that need special allocators, particularly for allocating out of some board-specific SRAM for DMA handles).

There are machvec functions added and removed over time, so always be sure to consult include/asm-sh/machvec.h for the current state of the machvec.

The kernel will automatically wrap in generic routines for undefined function pointers in the machvec at boot time, as machvec functions are referenced unconditionally throughout most of the tree. Some boards have incredibly sparse machvecs (such as the dreamcast and sh03), whereas others must define virtually everything (rts7751r2d).

Adding a new machine is relatively trivial (using vapor as an example):

If the board-specific definitions are quite minimalistic, as is the case for the vast majority of boards, simply having a single board-specific header is sufficient.

- add a new file include/asm-sh/vapor.h which contains prototypes for any machine specific IO functions prefixed with the machine name, for example vapor_inb. These will be needed when filling out the machine vector.

Note that these prototypes are generated automatically by setting `__IO_PREFIX` to something sensible. A typical example would be:

```
#define __IO_PREFIX vapor
#include <asm/io_generic.h>
```

somewhere in the board-specific header. Any boards being ported that still have a legacy io.h should remove it entirely and switch to the new model.

- Add machine vector definitions to the board’s setup.c. At a bare minimum, this must be defined as something like:

```
struct sh_machine_vector mv_vapor __initmv = {
    .mv_name = "vapor",
};
ALIAS_MV(vapor)
```

- finally add a file arch/sh/boards/vapor/io.c, which contains definitions of the machine specific io functions (if there are enough to warrant it).

2.3 3. Hooking into the Build System

Now that we have the corresponding directories setup, and all of the board-specific code is in place, it's time to look at how to get the whole mess to fit into the build system.

Large portions of the build system are now entirely dynamic, and merely require the proper entry here and there in order to get things done.

The first thing to do is to add an entry to arch/sh/Kconfig, under the "System type" menu:

```
config SH_VAPOR
    bool "Vapor"
    help
    select Vapor if configuring for a FooTech Vaporboard.
```

next, this has to be added into arch/sh/Makefile. All boards require a machdir-y entry in order to be built. This entry needs to be the name of the board directory as it appears in arch/sh/boards, even if it is in a sub-directory (in which case, all parent directories below arch/sh/boards/ need to be listed). For our new board, this entry can look like:

```
machdir-$(CONFIG_SH_VAPOR) += vapor
```

provided that we've placed everything in the arch/sh/boards/vapor/ directory.

Next, the build system assumes that your include/asm-sh directory will also be named the same. If this is not the case (as is the case with multiple boards belonging to a common family), then the directory name needs to be implicitly appended to incdir-y. The existing code manages this for the Solution Engine and hp6xx boards, so see these for an example.

Once that is taken care of, it's time to add an entry for the mach type. This is done by adding an entry to the end of the arch/sh/tools/mach-types list. The method for doing this is self explanatory, and so we won't waste space restating it here. After this is done, you will be able to use implicit checks for your board if you need this somewhere throughout the common code, such as:

```
/* Make sure we're on the FooTech Vaporboard */
if (!mach_is_vapor())
    return -ENODEV;
```

also note that the mach_is_boardname() check will be implicitly forced to lower-case, regardless of the fact that the mach-types entries are all uppercase. You can read the script if you really care, but it's pretty ugly, so you probably don't want to do that.

Now all that's left to do is providing a defconfig for your new board. This way, other people who end up with this board can simply use this config for reference instead of trying to guess what settings are supposed to be used on it.

Also, as soon as you have copied over a sample .config for your new board (assume arch/sh/configs/vapor_defconfig), you can also use this directly as a build target, and it will be implicitly listed as such in the help text.

Looking at the ‘make help’ output, you should now see something like:

Architecture specific targets (sh):

zImage	Compressed (arch/sh/boot/zImage)	kernel	image
adx_defconfig	Build for adx		
cqreek_defconfig	Build for cqreek		
dream- cast_defconfig	Build for dreamcast		
...			
vapor_defconfig	Build for vapor		

which then allows you to do:

```
$ make ARCH=sh CROSS_COMPILE=sh4-linux- vapor_defconfig vmlinux
```

which will in turn copy the defconfig for this board, run it through oldconfig (prompting you for any new options since the time of creation), and start you on your way to having a functional kernel for your new board.

NOTES ON REGISTER BANK USAGE IN THE KERNEL

3.1 Introduction

The SH-3 and SH-4 CPU families traditionally include a single partial register bank (selected by SR.RB, only r0 ...r7 are banked), whereas other families may have more full-featured banking or simply no such capabilities at all.

3.2 SR.RB banking

In the case of this type of banking, banked registers are mapped directly to r0 ...r7 if SR.RB is set to the bank we are interested in, otherwise ldc/stc can still be used to reference the banked registers (as r0_bank ...r7_bank) when in the context of another bank. The developer must keep the SR.RB value in mind when writing code that utilizes these banked registers, for obvious reasons. Userspace is also not able to poke at the bank1 values, so these can be used rather effectively as scratch registers by the kernel.

Presently the kernel uses several of these registers.

- r0_bank, r1_bank (referenced as k0 and k1, used for scratch registers when doing exception handling).
- r2_bank (used to track the EXPEVT/INTEVT code)
 - Used by do_IRQ() and friends for doing irq mapping based off of the interrupt exception vector jump table offset
- r6_bank (global interrupt mask)
 - The SR.IMASK interrupt handler makes use of this to set the interrupt priority level (used by local_irq_enable())
- r7_bank (current)

MEMORY MANAGEMENT

4.1 SH-4

4.1.1 Store Queue API

void **sq_flush_range**(unsigned long start, unsigned int len)

Flush (prefetch) a specific SQ range

Parameters

unsigned long start

the store queue address to start flushing from

unsigned int len

the length to flush

Description

Flushes the store queue cache from **start** to **start + len** in a linear fashion.

unsigned long **sq_remap**(unsigned long phys, unsigned int size, const char
*name, pgprot_t prot)

Map a physical address through the Store Queues

Parameters

unsigned long phys

Physical address of mapping.

unsigned int size

Length of mapping.

const char *name

User invoking mapping.

pgprot_t prot

Protection bits.

Description

Remaps the physical address **phys** through the next available store queue address of **size** length. **name** is logged at boot time as well as through the sysfs interface.

void **sq_unmap**(unsigned long vaddr)
Unmap a Store Queue allocation

Parameters

unsigned long vaddr
Pre-allocated Store Queue mapping.

Description

Unmaps the store queue allocation **map** that was previously created by [*sq_remap\(\)*](#). Also frees up the pte that was previously inserted into the kernel page table and discards the UTLB translation.

MACHINE SPECIFIC INTERFACES

5.1 mach-dreamcast

int **aica_rtc_gettimeofday**(struct device *dev, struct rtc_time *tm)

Get the time from the AICA RTC

Parameters

struct device *dev

the RTC device (ignored)

struct rtc_time *tm

pointer to resulting RTC time structure

Description

Grabs the current RTC seconds counter and adjusts it to the Unix Epoch.

int **aica_rtc_settimeofday**(struct device *dev, struct rtc_time *tm)

Set the AICA RTC to the current time

Parameters

struct device *dev

the RTC device (ignored)

struct rtc_time *tm

pointer to new RTC time structure

Description

Adjusts the given **tv** to the AICA Epoch and sets the RTC seconds counter.

5.2 mach-x3proto

int **ilsel_enable**(ilsel_source_t set)

Enable an ILSEL set.

Parameters

ilsel_source_t set

ILSEL source (see `ilsel_source_t` enum in `include/asm-sh/ilsel.h`).

Description

Enables a given non-aliased ILSEL source (\leq ILSEL_KEY) at the highest available interrupt level. Callers should take care to order callsites noting descending interrupt levels. Aliasing FPGA and external board IRQs need to use *ilssel_enable_fixed()*.

The return value is an IRQ number that can later be taken down with *ilssel_disable()*.

int **ilssel_enable_fixed**(ilssel_source_t set, unsigned int level)

Enable an ILSEL set at a fixed interrupt level

Parameters

ilssel_source_t set

ILSEL source (see ilssel_source_t enum in include/asm-sh/ilssel.h).

unsigned int level

Interrupt level (1 - 15)

Description

Enables a given ILSEL source at a fixed interrupt level. Necessary both for level reservation as well as for aliased sources that only exist on special ILSEL#s.

Returns an IRQ number (as *ilssel_enable()*).

void **ilssel_disable**(unsigned int irq)

Disable an ILSEL set

Parameters

unsigned int irq

Bit position for ILSEL set value (retval from enable routines)

Description

Disable a previously enabled ILSEL set.

6.1 SuperHyway

```
int superhyway_add_device(unsigned long base, struct superhyway_device  
                        *sdev, struct superhyway_bus *bus)
```

Add a SuperHyway module

Parameters

unsigned long base

Physical address where module is mapped.

struct superhyway_device *sdev

SuperHyway device to add, or NULL to allocate a new one.

struct superhyway_bus *bus

Bus where SuperHyway module resides.

Description

This is responsible for adding a new SuperHyway module. This sets up a new struct superhyway_device for the module being added if **sdev** == NULL.

Devices are initially added in the order that they are scanned (from the top-down of the memory map), and are assigned an ID based on the order that they are added. Any manual addition of a module will thus get the ID after the devices already discovered regardless of where it resides in memory.

Further work can and should be done in superhyway_scan_bus(), to be sure that any new modules are properly discovered and subsequently registered.

```
int superhyway_register_driver(struct superhyway_driver *drv)
```

Register a new SuperHyway driver

Parameters

struct superhyway_driver *drv

SuperHyway driver to register.

Description

This registers the passed in **drv**. Any devices matching the id table will automatically be populated and handed off to the driver's specified probe routine.

void **superhyway_unregister_driver**(struct superhyway_driver *drv)

Unregister a SuperHyway driver

Parameters

struct superhyway_driver *drv

SuperHyway driver to unregister.

Description

This cleans up after [superhyway_register_driver\(\)](#), and should be invoked in the exit path of any module drivers.

6.2 Maple

int **maple_driver_register**(struct maple_driver *drv)

register a maple driver

Parameters

struct maple_driver *drv

maple driver to be registered.

Description

Registers the passed in **drv**, while updating the bus type. Devices with matching function IDs will be automatically probed.

void **maple_driver_unregister**(struct maple_driver *drv)

unregister a maple driver.

Parameters

struct maple_driver *drv

maple driver to unregister.

Description

Cleans up after [maple_driver_register\(\)](#). To be invoked in the exit path of any module drivers.

void **maple_getcond_callback**(struct maple_device *dev, void (*callback)(struct mapleq *mq), unsigned long interval, unsigned long function)

setup handling MAPLE_COMMAND_GETCOND

Parameters

struct maple_device *dev

device responding

void (*callback) (struct mapleq *mq)

handler callback

unsigned long interval

interval in jiffies between callbacks

unsigned long function

the function code for the device

int **maple_add_packet**(struct maple_device *mdev, u32 function, u32 command,
size_t length, void *data)

add a single instruction to the maple bus queue

Parameters

struct maple_device *mdev

maple device

u32 function

function on device being queried

u32 command

maple command to add

size_t length

length of command string (in 32 bit words)

void *data

remainder of command string

INDEX

`\spxentryaica_rtc_gettimeofday\spxextraC`
function, [15](#)

`\spxentryaica_rtc_settimeofday\spxextraC`
function, [15](#)

`\spxentryilssel_disable\spxextraC` func-
tion, [16](#)

`\spxentryilssel_enable\spxextraC` func-
tion, [15](#)

`\spxentryilssel_enable_fixed\spxextraC`
function, [16](#)

`\spxentrymaple_add_packet\spxextraC`
function, [19](#)

`\spxentrymaple_driver_register\spxextraC`
function, [18](#)

`\spxentrymaple_driver_unregister\spxextraC`
function, [18](#)

`\spxentrymaple_getcond_callback\spxextraC`
function, [18](#)

`\spxentrysq_flush_range\spxextraC`
function, [13](#)

`\spxentrysq_remap\spxextraC` function,
[13](#)

`\spxentrysq_unmap\spxextraC` function,
[13](#)

`\spxentrysuperhyway_add_device\spxextraC`
function, [17](#)

`\spxentrysuperhyway_register_driver\spxextraC`
function, [17](#)

`\spxentrysuperhyway_unregister_driver\spxextraC`
function, [17](#)